

# ICS 663: Homework 3 - Linear Discriminant Functions

Christopher Mullins

November 11, 2011

---

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Experiments</b>	<b>2</b>
2.1	Perceptron . . . . .	2
2.2	MSE . . . . .	3
<b>3</b>	<b>Implementation</b>	<b>3</b>

---

## 1 Introduction

Suppose that a dataset  $D = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$  has two classes, and that  $d_i$  denotes the class for  $\mathbf{x}_i$ . For convenience, we can label these classes  $-1$  and  $1$ , so that if  $\mathbf{x}_i$  is in the first class,  $d_i = -1$ .

A dataset with two classes is linearly separable if there exists a vector  $\mathbf{a}$  such that  $d_i \cdot \mathbf{a}^T \mathbf{x}_i > 0 \forall i$ . It is not uncommon for datasets to be *approximately* linearly separable, meaning that only a minority of data are misclassified by an appropriately chosen  $\mathbf{a}$ . For this reason, methods for finding linear separators are of some interest.

The perceptron is a very simple artificial neural network that is capable of finding a linear separator for a dataset by incrementally updating  $\mathbf{a}$  based on misclassified results. More explicitly, at each iteration  $i$  and for each misclassified sample  $\mathbf{x}_j$ , we update  $\mathbf{a}$  using the following rule:

$$\mathbf{a}_{i+1} = \mathbf{a}_i + \eta(i) (d_j - \mathbf{a}_i^T \mathbf{x}_j) \mathbf{x}_j.$$

Here,  $\eta(i)$  denotes the *learning rate* at iteration  $i$ . This is generally a value between 0 and 2, and can change (usually decay) as  $i$  gets larger. The perceptron halts when no misclassified samples remain. For data that aren't fully linearly separable, a different halting condition should be used (error threshold, number of iterations, etc.).

The MSE method uses a well-known method for linear regression to find a linear boundary. We first form a matrix with the data:

$$\mathbf{Y} = \begin{pmatrix} d_0 x_{0,0} & d_0 x_{0,1} & \dots & d_0 x_{0,k} & d_0 \\ d_1 x_{1,0} & d_1 x_{1,1} & \dots & d_1 x_{1,k} & d_1 \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ d_n x_{n,0} & d_n x_{n,1} & \dots & d_n x_{n,k} & d_n \end{pmatrix}.$$

We then invent (or calculate) some margin weight vector  $\mathbf{b}$  (such that  $b_i > 0 \forall i$ ) that we solve for, so that  $\mathbf{Y}\mathbf{a} = \mathbf{b}$ . Since  $\mathbf{Y}$  is almost never square in practice ( $n \gg k$  in general), the system is *overdetermined*, meaning there is likely not an exact solution. We can, however, find a solution that minimizes the squared error,  $\|\mathbf{Y}\mathbf{a} - \mathbf{b}\|$ .

We can use the pseudoinverse to find such a solution:

$$\mathbf{Y}^\dagger = (\mathbf{Y}^T \mathbf{Y})^{-1} \mathbf{Y}^T.$$

This calculation works if  $\mathbf{Y}^T \mathbf{Y}$  is non-singular. If it is singular, one can use the more generalized definition or use a gradient descent method. More generally, the pseudoinverse is:

$$\lim_{\epsilon \rightarrow 0} (\mathbf{Y}^T \mathbf{Y} + \epsilon \mathbf{I})^{-1} \mathbf{Y}^T.$$

## 2 Experiments

For this assignment, I use the perceptron and MSE methods for finding linear discriminant functions. The provided data is shown in figure 1. In the result plots, I omit the legend to leave more room for the separators. The same color and symbols for each class are used in these plots as in figure 1.




Figure 1: Data provided for homework assignment

### 2.1 Perceptron

With the data provided, the pseudoinverse  $\mathbf{Y}^\dagger$  exists, so using gradient descent is unnecessary.

I ran my perceptron code on the provided data four times and provided the results in a single plot. The black lines represent the decision boundary  $\mathbf{a}$  at each epoch. The green line represents the final result produced by the perceptron. It perfectly separates the data into two classes. The results are shown in figure 2 and 3 for  $\eta = 0.5$  and  $\eta = 0.01$ , respectively.

The results are unsurprising. With  $\eta = 0.5$ , a decision boundary is found within a few iterations. With  $\eta = 0.01$ , significantly more iterations are needed before the perceptron converges on a solution.

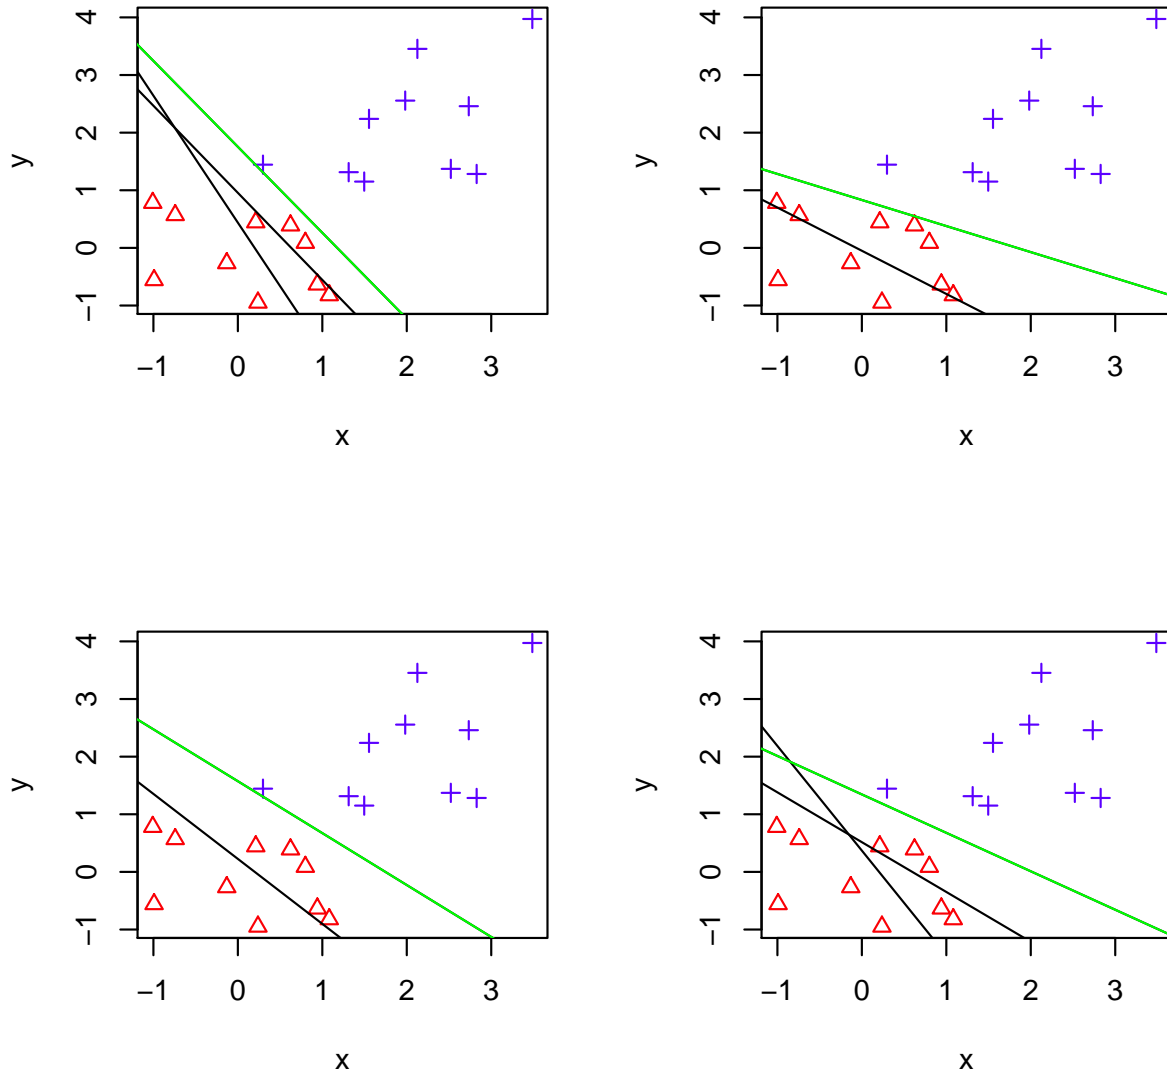


Figure 2: Results for Perceptron with  $\eta = 0.5$ .

## 2.2 MSE

One disadvantage of the MSE method for finding a linear decision boundary is that it does not guarantee a perfect separator, even if the data is perfectly separable. In the provided data, the solution provided by MSE does *not* perfectly separate the data. The results are shown in figure 4.

## 3 Implementation

The implementation I used for the plots included in this report is done in R. The code is split across three files: `perceptron.R`, `mse.R` and `hw3.R`. The former two have code for the perceptron and MSE methods, respectively. The last is a script that seamlessly runs all of the tasks required for HW3 (generates plots and

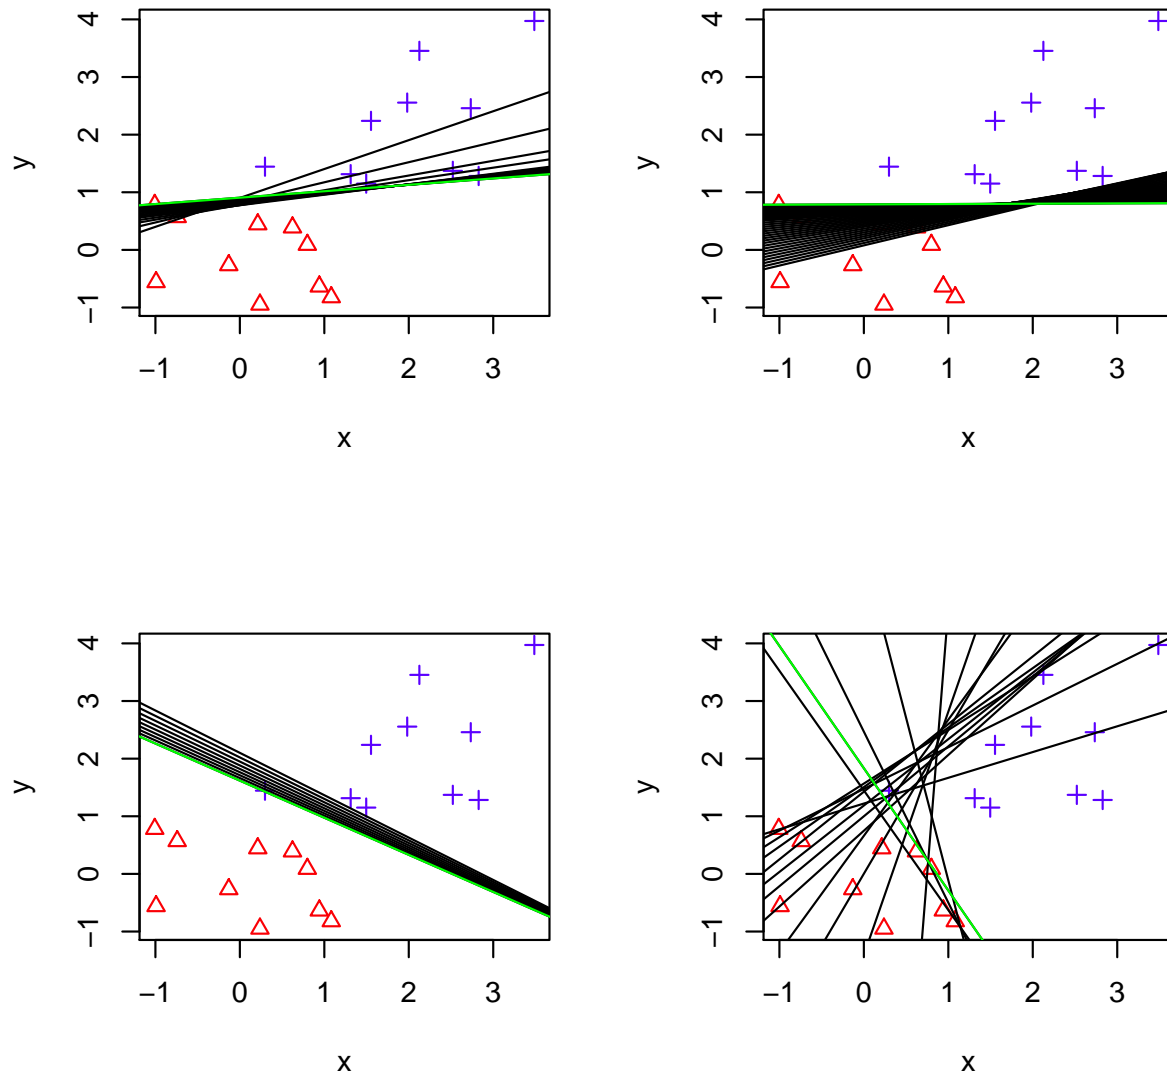


Figure 3: Results for Perceptron with  $\eta = 0.01$ .

puts them in the `./plots` directory).

If you'd like to run the code, you can use the command `R --no-save --slave < hw3.R`.




Figure 4: Results for MSE